



**Big Data & Analytics**  
**(24ECSC401)**

Course Project Report On

**Twitter sentiment analysis using Hadoop**

*submitted in partial fulfillment of the requirement for the degree of*

Bachelor of Engineering

In

School of Computer Science and Engineering

Submitted By

<b>Prasad Palled</b>	<b>01FE21BCS123</b>	<b>119</b>
<b>Saisamarth Udikeri</b>	<b>01FE21BCS003</b>	<b>167</b>
<b>Rishi Hiremath</b>	<b>01FE21BCS304</b>	<b>150</b>
<b>Konkathi Rithin Kumar</b>	<b>01FE20BCS008</b>	<b>162</b>

**Under the guidance of**

Dr. Suvarna Kanakraddi

**SCHOOL OF COMPUTER SCIENCE & ENGINEERING**

**HUBLI-580 031 (India).**

**Academic year 2024-25**

# **TABLE OF CONTENT**

## **Chapters**

- 1. Introduction**
- 2. Problem Statement**
- 3. Data set description**
- 4. Input data analysis using big data sandbox tool**
- 5. Model Selection**
- 6. Hadoop map reduce implementation in Java.**
- 7. Result Analysis in Hadoop**
- 8. Conclusion and Future Scope**
- 9. Reference**

# Chapter 1

## INTRODUCTION

In the era of social media dominance, platforms like Twitter serve as a dynamic repository of real public opinions, emotions, and reactions. Analysing this massive influx of data offers valuable insights for businesses, governments, and researchers alike. However, the real challenge lies in processing this data efficiently to derive actionable insights in real time. This project, Sentiment Analysis, addresses this challenge by leveraging Hadoop's MapReduce framework to classify tweets into positive, negative, or neutral sentiments as they are posted.

The system employs a hybrid model that combines collaborative filtering and content-based recommendation techniques to enhance both the accuracy and scalability of the sentiment analysis process. The data processing is managed entirely through Hadoop, with Mapper and Reducer tasks designed to process the large-scale unstructured data efficiently. This approach ensures seamless handling of the tweet data in a distributed environment, making the analysis both scalable and robust.

The significance of this project lies in its ability to transform raw social media data into meaningful insights at scale. By utilizing Hadoop's distributed computing capabilities, the system provides a competitive edge to stakeholders by enabling timely decision-making and a deeper understanding of public sentiment. This work demonstrates the potential of big data technologies to revolutionize sentiment analysis and set the stage for future advancements in real analytics systems.

# Chapter 2

## PROBLEM STATEMENT

The project addresses the challenge of analysing Twitter's fast-paced, high-volume data by building a sentiment analysis system using Hadoop's MapReduce framework for efficient and accurate processing of large datasets.

1. Twitter generates massive, fast-paced data streams, making it challenging to analyse and derive insights efficiently.
2. Traditional sentiment analysis methods are inefficient at handling the scale, speed, and structure of Twitter data.
3. There is a need for a robust, scalable system that uses MapReduce to process and classify sentiments accurately, leveraging machine learning models trained on large datasets.

### Objectives

1. Develop a scalable system using Hadoop's MapReduce framework to process and analyze Twitter data efficiently.
2. Implement machine learning algorithms trained via MapReduce for sentiment classification.
3. Validate the sentiment analysis model using a validation mapper for accurate results.
4. Provide actionable insights for timely and informed decision-making.

### Technologies and Tools

- **Hadoop with MapReduce:** For distributed training, validation, and analysis of datasets.
- **Datasets:** For training and validating the sentiment analysis model.
- **Java:** For implementing MapReduce workflows and machine learning models.

# Chapter 3

## DATASET DESCRIPTION

This dataset contains 52,24,794 tweets with metadata and analysis outcomes across 20 columns. Key features include tweet IDs, timestamps, text, language flags, sentiment and emotion labels, engagement metrics, user activity, and advanced features like code-switching and automation detection. Some columns have missing data, such as country-specific sentiment. It supports sentiment analysis, language detection, and user behavior studies.

### General Details:

- **Total Records:** 52,24,794 rows
- **Total Columns:** 20
- **Data Types:** Includes integer, float, and string data types.

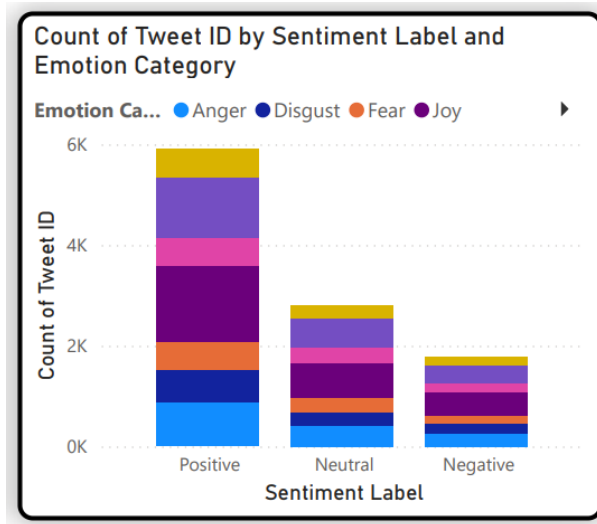
### Column Description:

Tweet ID	Unique identifier for each tweet.
Country	Country where the tweet was posted.
Date	Date when the tweet was posted.
Tweet	Content of the tweet.
Definitely English	Indicates if the tweet is classified as definitely in English (1 = Yes, 0 = No).
Ambiguous	Indicates if the language is ambiguous (1 = Yes, 0 = No).
Definitely Not English	Indicates if the tweet is classified as not in English (1 = Yes, 0 = No).
Code-Switched	Indicates if the tweet contains multiple languages (1 = Yes, 0 = No).
Ambiguous due to Named Entities	Indicates if ambiguity arises from named entities (1 = Yes, 0 = No).
Automatically Generated Tweets	Indicates if the tweet was auto-generated (1 = Yes, 0 = No).
Emotion Category	Emotion expressed in the tweet (e.g., Joy, Fear, Disgust).
Sentiment Label	Overall sentiment classification (Positive, Negative, Neutral).
Sentiment by Country	Sentiment specific to the country context.
Likes	Number of likes received by the tweet.
Retweets	Number of times the tweet was retweeted.
Replies	Number of replies to the tweet.
Tweet Length	Length of the tweet in characters.
User Activity Level	Classification of the user's activity level (e.g., Highly Active).
Is Automated	Indicates if the tweet was posted by an automated system (Yes or No).
Code-Switched Sentiment	Sentiment analysis for code-switched tweets (if applicable).

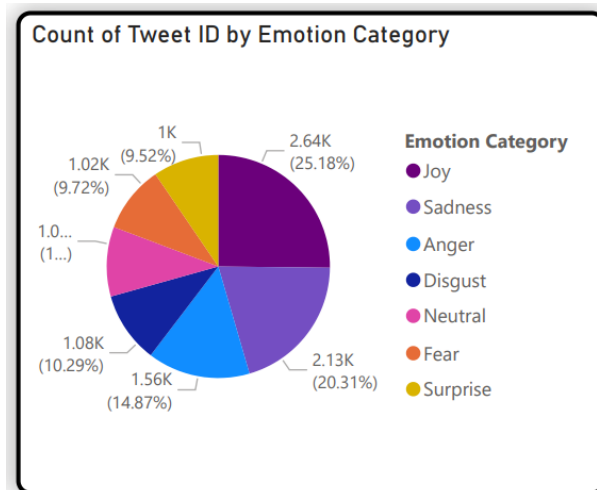
## Chapter 4

# INPUT DATA ANALYSIS USING BIG DATA SANDBOX TOOL

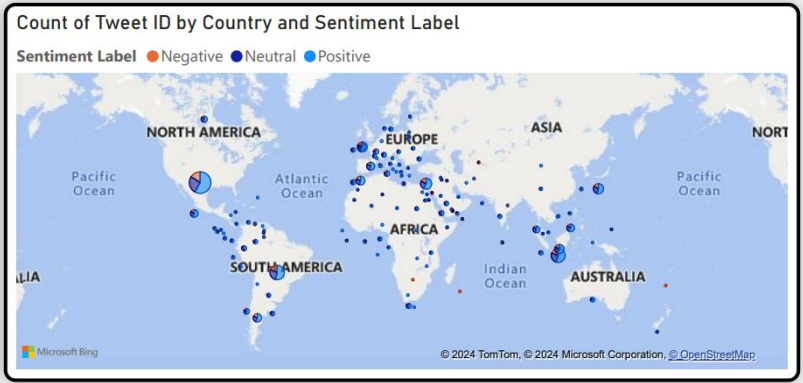
1. The bar graph shows the count of tweets categorized by Sentiment Label (e.g., Positive, Negative) and Emotion Category (e.g., Joy, Fear). Each bar represents the number of tweets within a specific sentiment and emotion combination. This visualization highlights the distribution of emotions across different sentiment labels, making it easy to identify dominant patterns and trends in the dataset.




2. The pie chart represents the distribution of tweets by **Emotion Category** (e.g., Joy, Fear, Disgust). Each segment reflects the proportion of tweets associated with a specific emotion, providing a clear overview of the emotional makeup of the dataset. This visualization helps identify the most and least prevalent emotions among the tweets.



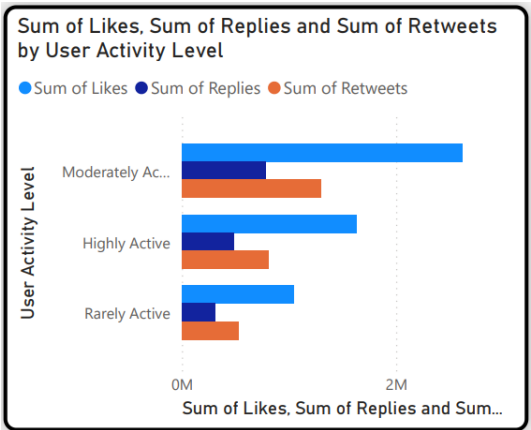
- The world map visualization displays the count of tweets categorized by **Country** and **Sentiment Label** (e.g., Positive, Negative). Each country is shaded or marked based on the number of tweets and their associated sentiment, providing a global view of sentiment distribution.



- The donut graph shows the distribution of tweets by **Country**, with each segment representing the proportion of tweets originating from a specific country. This visualization provides a clear and concise overview of the geographical spread.



- The **horizontal stacked bar chart** shows the **Sum of Likes, Replies, and Retweets** for each **User Activity Level**. It allows for easy comparison of engagement metrics across different levels of user activity.



# Chapter 5

## MODEL SELECTION

Model selection is a crucial phase in the machine learning process, focused on identifying the most suitable algorithm or model to address a specific problem effectively. It involves comparing multiple models and choosing one that delivers optimal performance while maintaining efficiency and generalizability. The objective is to strike a balance between predictive accuracy, computational cost, and the interpretability of results.

The first step in model selection is understanding the problem type—whether it's classification, regression, clustering, or another machine learning task. For example, logistic regression is a common choice for binary classification problems, while decision trees or ensemble models like random forests may perform better for datasets with complex relationships. Similarly, for deep insights into large-scale datasets, neural networks or deep learning approaches might be more appropriate.

### **Key factors that influence model selection include:**

1. **Performance Metrics:** Different problems demand different evaluation metrics. Classification models are assessed using metrics such as accuracy, precision, recall, and F1-score, while regression models rely on metrics like mean squared error or R-squared. These metrics help quantify the model's effectiveness.
2. **Model Complexity:** A simpler model is generally preferred when it achieves comparable performance. Simple models reduce the risk of overfitting and are easier to interpret. For instance, linear regression may suffice for linearly separable data, but more complex models like support vector machines or ensemble techniques are better suited for non-linear patterns.
3. **Scalability:** Models need to handle the volume and complexity of data effectively. Scalable algorithms like distributed implementations of decision trees or gradient-boosted methods are crucial in big data scenarios.
4. **Cross-Validation:** Robust techniques like k-fold cross-validation are essential to evaluate a model's performance on unseen data and prevent overfitting. These techniques ensure that the model generalizes well across different subsets of the data.

Tools such as grid search and automated machine learning frameworks can further aid in the model selection process by tuning hyperparameters and systematically comparing models. Domain knowledge also plays a significant role, as understanding the data and its context helps narrow down the most appropriate models.

Ultimately, model selection is about finding the best trade-off between accuracy, efficiency, interpretability, and robustness. It ensures the chosen model aligns with the project's goals, delivering actionable insights and reliable predictions in real-world applications.



# Chapter 6

## HADOOP MAP REDUCE IMPLEMENTATION IN JAVA

### 1. Sentiment Analysis Operation:

#### Description:

This operation involves analyzing textual data to classify its sentiment into categories such as Positive, Negative, or Neutral. It includes modules for training sentiment analysis models, validating their accuracy, and orchestrating the process using a driver class. The Mapper and Reducer components process and aggregate data to prepare it for sentiment-based insights.

#### 1) Training Mapper:

The Sentimental Training Mapper Java code is designed to process and prepare data for sentiment analysis training. It maps input data, such as text or metadata, to corresponding sentiment labels (e.g., Positive, Negative). The code typically includes logic for parsing, filtering, and transforming the data into a format suitable for training sentiment analysis models. It serves as a critical component in building and optimizing machine learning models for sentiment classification tasks.

#### Code:

```
1  import org.apache.hadoop.io.LongWritable;
2  import org.apache.hadoop.io.Text;
3  import org.apache.hadoop.mapreduce.Mapper;
4
5  import java.io.IOException;
6
7  public class TrainingMapper extends Mapper<LongWritable, Text, Text, Text> {
8
9      @Override
10     protected void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
11         String line = value.toString();
12         String[] fields = line.split(regex:"(?=<[^\"]*\"[^\"]*\">)*[^\"]*$"); // Handle quoted CSV fields
13
14         if (fields.length > 10) { // Validate minimum fields (Tweet + Sentiment)
15             String tweet = fields[3].toLowerCase().replaceAll(regex:"[a-zA-Z0-9#@ ]", replacement:""); // Clean tweet text
16             String sentiment = fields[11].trim().toLowerCase(); // Trim sentiment for safety
17
18             for (String word : tweet.split(regex:" ")) { // Tokenize tweet text by spaces
19                 word = word.trim(); // Remove leading/trailing spaces
20                 if (!word.isEmpty() && word.length() > 2) { // Skip empty and short words
21                     context.write(new Text(word), new Text(sentiment));
22                 }
23             }
24         } else {
25             // Log invalid row for debugging
26             System.err.println("Invalid row: " + line);
27         }
28     }
29 }
30
```

## 2) Training Reducer

The Sentimental Training Reducer is responsible for aggregating and summarizing data during the sentiment training process. It combines outputs from the Sentimental Training Mapper and processes them to generate consolidated results, such as sentiment label counts or patterns, that are essential for training sentiment analysis models.

### Code:

```
1  import org.apache.hadoop.io.Text;
2  import org.apache.hadoop.mapreduce.Reducer;
3
4  import java.io.IOException;
5
6  public class TrainingReducer extends Reducer<Text, Text, Text, Text> {
7
8      @Override
9      protected void reduce(Text key, Iterable<Text> values, Context context) throws IOException, InterruptedException {
10         int positive = 0, negative = 0, neutral = 0;
11
12         // Count sentiments
13         for (Text value : values) {
14             String sentiment = value.toString().toLowerCase();
15             switch (sentiment) {
16                 case "positive":
17                     positive++;
18                     break;
19                 case "negative":
20                     negative++;
21                     break;
22                 case "neutral":
23                     neutral++;
24                     break;
25             }
26         }
27
28         // Determine dominant sentiment
29         String dominantSentiment = "neutral";
30         if (positive > negative && positive > neutral) {
31             dominantSentiment = "positive";
32         } else if (negative > positive && negative > neutral) {
33             dominantSentiment = "negative";
34         }
35
36         // Emit word and its dominant sentiment only if counts are non-zero
37         if (positive + negative + neutral > 0) {
38             context.write(key, new Text(dominantSentiment));
39         } else {
40             // Log words with no sentiment (optional for debugging)
41             System.err.println("No sentiment found for word: " + key.toString());
42         }
43     }
44 }
45
```

### 3) Validation Mapper:

The Sentimental Validation Mapper processes input data for sentiment validation. It maps textual data or metadata to sentiment categories while ensuring the data adheres to the validation criteria. This is often used for testing or evaluating sentiment analysis models.

#### Code:

```
1 import org.apache.hadoop.io.LongWritable;
2 import org.apache.hadoop.io.Text;
3 import org.apache.hadoop.mapreduce.Mapper;
4 import org.apache.hadoop.fs.Path;
5 import org.apache.hadoop.fs.FileSystem;
6
7
8 public class ValidationMapper extends Mapper<LongWritable, Text, Text, Text> {
9     private Map<String, String> sentimentDictionary = new HashMap<>();
10
11     @Override
12     protected void setup(Context context) throws IOException {
13         URL[] cacheFiles = context.getCacheFiles(); // Load the sentiment dictionary from distributed cache
14         if (cacheFiles != null && cacheFiles.length > 0) {
15             FileSystem fs = FileSystem.get(context.getConfiguration());
16             BufferedReader reader = new BufferedReader(new InputStreamReader(fs.open(new Path(cacheFiles[0]))));
17             String line;
18             while ((line = reader.readLine()) != null) {
19                 String[] fields = line.split(regex:"\\t"); // Key-value pairs from training output
20                 if (fields.length == 2) {
21                     sentimentDictionary.put(fields[0], fields[1]);
22                 }
23             }
24             reader.close();
25         }
26     }
27
28     @Override
29     protected void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
30         String line = value.toString();
31         String[] fields = line.split(regex:","); // Split by comma for CSV
32
33         if (fields.length > 10) { // Ensure there are enough fields
34             String tweet = fields[3]; // Tweet text
35             String actualSentiment = fields[11]; // Sentiment Label column
36
37             int positive = 0, negative = 0, neutral = 0;
38
39             for (String word : tweet.split(regex:" ")) {
40                 word = word.replaceAll(regex:"[^a-zA-Z]", replacement:"").toLowerCase();
41                 String predictedSentiment = sentimentDictionary.get(word);
42                 if (predictedSentiment != null) {
43                     switch (predictedSentiment) {
44                         case "positive":
45                             positive++;
46                             break;
47                         case "negative":
48                             negative++;
49                             break;
50                         case "neutral":
51                             neutral++;
52                             break;
53                     }
54                 }
55             }
56
57             String predictedSentiment = "neutral";
58             if (positive > negative && positive > neutral) predictedSentiment = "positive";
59             if (negative > positive && negative > neutral) predictedSentiment = "negative";
60
61             // Emit the tweet text and its predicted sentiment
62             context.write(new Text(tweet), new Text(predictedSentiment));
63         }
64     }
65 }
```

#### 4) Driver Class:

The Sentimental Driver acts as the main driver class, coordinating the execution of the sentiment analysis process. It sets up the configuration for the Mapper and Reducer, manages the job flow, and oversees the entire process from input to output.

#### Code:

```
1 import org.apache.hadoop.conf.Configuration;
2 import org.apache.hadoop.fs.Path;
3 import org.apache.hadoop.io.Text;
4 import org.apache.hadoop.mapreduce.Job;
5 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
6 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
7
8 public class SentimentAnalysisDriver {
9     Run | Debug
10     public static void main(String[] args) throws Exception {
11         if (args.length < 3) {
12             System.err.println("Usage: SentimentAnalysisDriver <training input> <validation input> <output>");
13             System.exit(-1);
14         }
15
16         Configuration conf = new Configuration();
17
18         // Job 1: Training Phase
19         Job trainingJob = Job.getInstance(conf, "Training Phase");
20         trainingJob.setJarByClass(SentimentAnalysisDriver.class);
21
22         trainingJob.setMapperClass(TrainingMapper.class);
23         trainingJob.setReducerClass(TrainingReducer.class);
24
25         trainingJob.setOutputKeyClass(Text.class);
26         trainingJob.setOutputValueClass(Text.class);
27
28         FileInputFormat.addInputPath(trainingJob, new Path(args[0])); // Training input
29         Path trainingOutput = new Path("/user/hadoop/training_output");
30         FileOutputFormat.setOutputPath(trainingJob, trainingOutput);
31
32         System.out.println("Starting training phase...");
33         if (!trainingJob.waitForCompletion(true)) {
34             System.err.println("Training phase failed.");
35             System.exit(1);
36         }
37
38         // Job 2: Validation Phase
39         Job validationJob = Job.getInstance(conf, "Validation Phase");
40         validationJob.setJarByClass(SentimentAnalysisDriver.class);
41
42         validationJob.setMapperClass(ValidationMapper.class);
43
44         validationJob.setOutputKeyClass(Text.class);
45         validationJob.setOutputValueClass(Text.class);
46
47         FileInputFormat.addInputPath(validationJob, new Path(args[1])); // Validation input
48         FileOutputFormat.setOutputPath(validationJob, new Path(args[2])); // Validation output
49
50         // Add the training output to the distributed cache
51         validationJob.addCacheFile(new Path("/user/hadoop/training_output/part-r-00000").toUri());
52
53         System.out.println("Starting validation phase...");
54         if (!validationJob.waitForCompletion(true)) {
55             System.err.println("Validation phase failed.");
56             System.exit(1);
57         }
58
59         System.out.println("Sentiment Analysis completed successfully.");
60     }
61 }
```

## 2. Emotion Analysis Operation

### Description:

The emotion analysis operation focuses on categorizing text into emotion labels like Joy, Fear, or Disgust. It includes training mappers and reducers to process and summarize emotion data and validate classifications. A driver class ensures seamless coordination of the emotion analysis pipeline for extracting emotional patterns from text.

### 1) Emotion Valid Mapper:

The Emotion Valid Mapper maps input data to validate and categorize emotions (e.g., Joy, Fear). It ensures the data is correctly formatted and classified based on predefined emotion validation logic, supporting emotion analysis tasks.

### Code:

```
1  import org.apache.hadoop.io.LongWritable;
2  import org.apache.hadoop.io.Text;
3  import org.apache.hadoop.mapreduce.Mapper;
4  import org.apache.hadoop.fs.Path;
5  import org.apache.hadoop.fs.FileSystem;
6
7  import java.io.BufferedReader;
8  import java.io.InputStreamReader;
9  import java.net.URI;
10 import java.io.IOException;
11 import java.util.HashMap;
12 import java.util.Map;
13
14 public class EmotionValidationMapper extends Mapper<LongWritable, Text, Text, Text> {
15     private Map<String, String> emotionDictionary = new HashMap<>();
16
17     @Override
18     protected void setup(Context context) throws IOException {
19         URI[] cacheFiles = context.getCacheFiles(); // Load the emotion dictionary from distributed cache
20         if (cacheFiles != null && cacheFiles.length > 0) {
21             FileSystem fs = FileSystem.get(context.getConfiguration());
22             BufferedReader reader = new BufferedReader(new InputStreamReader(fs.open(new Path(cacheFiles[0]))));
23             String line;
24             while ((line = reader.readLine()) != null) {
25                 String[] fields = line.split(regex:"\\t"); // Key-value pairs from training output
26                 if (fields.length == 2) {
27                     emotionDictionary.put(fields[0], fields[1]);
28                 }
29             }
30             reader.close();
31         }
32     }
33
34     @Override
35     protected void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
36         String line = value.toString();
37         String[] fields = line.split(regex:","); // Split by comma for CSV
38
39         if (fields.length > 10) { // Ensure there are enough fields
40             String tweet = fields[3]; // Tweet text
41             String actualEmotion = fields[10]; // Emotion Label column
42
43             Map<String, Integer> emotionCounts = new HashMap<>();
44
45             for (String word : tweet.split(regex:" ")) {
46                 word = word.replaceAll(regex:"[a-zA-Z]", replacement:"").toLowerCase();
47                 String predictedEmotion = emotionDictionary.get(word);
48                 if (predictedEmotion != null) {
49                     emotionCounts.put(predictedEmotion, emotionCounts.getOrDefault(predictedEmotion, default:0) + 1);
50                 }
51             }
52
53             String predictedEmotion = emotionCounts.entrySet().stream()
54                 .max(Map.Entry.comparingByValue())
55                 .map(Map.Entry::getKey)
56                 .orElse(other:"neutral");
57
58             // Emit the tweet text and its predicted emotion
59             context.write(new Text(tweet), new Text(predictedEmotion));
60         }
61     }
62 }
63
```

## 2) Emotion Training Reducer:

The Emotion Training Reducer consolidates and processes outputs from the Emotion Training Mapper, aggregating emotion-related data. It helps generate summarized information, such as emotion frequencies or trends, for training emotion analysis models.

### Code:

```
1 import org.apache.hadoop.io.Text;
2 import org.apache.hadoop.mapreduce.Reducer;
3
4 import java.io.IOException;
5
6 public class EmotionTrainingReducer extends Reducer<Text, Text, Text, Text> {
7
8     @Override
9     protected void reduce(Text key, Iterable<Text> values, Context context) throws IOException, InterruptedException {
10         int joy = 0, sadness = 0, anger = 0, disgust = 0, neutral = 0, fear = 0, surprise = 0;
11
12         // Count emotions
13         for (Text value : values) {
14             String emotion = value.toString().toLowerCase();
15             switch (emotion) {
16                 case "joy":
17                     joy++;
18                     break;
19                 case "sadness":
20                     sadness++;
21                     break;
22                 case "anger":
23                     anger++;
24                     break;
25                 case "disgust":
26                     disgust++;
27                     break;
28                 case "neutral":
29                     neutral++;
30                     break;
31                 case "fear":
32                     fear++;
33                     break;
34                 case "surprise":
35                     surprise++;
36                     break;
37             }
38         }
39
40         // Determine dominant emotion
41         String dominantEmotion = "neutral"; // Default
42         int maxCount = Math.max(Math.max(joy, sadness), Math.max(anger, Math.max(disgust, Math.max(fear, surprise))));
43         if (joy == maxCount) dominantEmotion = "joy";
44         else if (sadness == maxCount) dominantEmotion = "sadness";
45         else if (anger == maxCount) dominantEmotion = "anger";
46         else if (disgust == maxCount) dominantEmotion = "disgust";
47         else if (fear == maxCount) dominantEmotion = "fear";
48         else if (surprise == maxCount) dominantEmotion = "surprise";
49
50         // Emit word and its dominant emotion
51         context.write(key, new Text(dominantEmotion));
52     }
53 }
54
```

### 3. Emotion Training Mapper:

The Emotion Training Mapper processes input data and maps it to specific emotion categories for training purposes. It performs initial transformations and classifications, serving as the first step in the emotion analysis training pipeline.

#### Code:

```
1  import org.apache.hadoop.io.LongWritable;
2  import org.apache.hadoop.io.Text;
3  import org.apache.hadoop.mapreduce.Mapper;
4
5  import java.io.IOException;
6
7  public class EmotionTrainingMapper extends Mapper<LongWritable, Text, Text, Text> {
8
9      @Override
10     protected void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
11         String line = value.toString();
12         String[] fields = line.split(regex:"(?:([^\"]|\"[^\"]*\")*(\"|$))*"); // Handle quoted CSV fields
13
14         if (fields.length > 10) { // Validate minimum fields (Tweet + Emotion)
15             String tweet = fields[3].toLowerCase().replaceAll(regex:"[^a-zA-Z0-9#@ ]", replacement:""); // Clean tweet text
16             String emotion = fields[10].trim().toLowerCase(); // Get emotion from field 10
17
18             for (String word : tweet.split(regex:" ")) { // Tokenize tweet text by spaces
19                 word = word.trim(); // Remove leading/trailing spaces
20                 if (!word.isEmpty() && word.length() > 2) { // Skip empty and short words
21                     context.write(new Text(word), new Text(emotion));
22                 }
23             }
24         } else {
25             // Log invalid row for debugging
26             System.err.println("Invalid row: " + line);
27         }
28     }
29 }
30
```



## 4. Emotion Analysis Driver:

The Emotion Analysis Driver orchestrates the emotion analysis job. It configures and manages the Mapper, Reducer, and input/output settings, ensuring smooth execution of the emotion analysis workflow from start to finish.

### Code:

```
1  // Source code is decompiled from a .class file using FernFlower decompiler.
2  import org.apache.hadoop.conf.Configuration;
3  import org.apache.hadoop.fs.Path;
4  import org.apache.hadoop.io.Text;
5  import org.apache.hadoop.mapreduce.Job;
6  import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
7  import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
8
9  public class EmotionAnalysisDriver {
10     public EmotionAnalysisDriver() {
11     }
12
13     public static void main(String[] var0) throws Exception {
14         if (var0.length < 3) {
15             System.err.println("Usage: EmotionAnalysisDriver <training input> <validation input> <output>");
16             System.exit(-1);
17         }
18
19         Configuration var1 = new Configuration();
20         Job var2 = Job.getInstance(var1, "Emotion Training Phase");
21         var2.setJarByClass(EmotionAnalysisDriver.class);
22         var2.setMapperClass(EmotionTrainingMapper.class);
23         var2.setReducerClass(EmotionTrainingReducer.class);
24         var2.setOutputKeyClass(Text.class);
25         var2.setOutputValueClass(Text.class);
26         FileInputFormat.addInputPath(var2, new Path(var0[0]));
27         Path var3 = new Path("/user/hadoop/emotion_training_output");
28         FileOutputFormat.setOutputPath(var2, var3);
29         System.out.println("Starting training phase...");
30         if (!var2.waitForCompletion(true)) {
31             System.err.println("Training phase failed.");
32             System.exit(1);
33         }
34
35         Job var4 = Job.getInstance(var1, "Emotion Validation Phase");
36         var4.setJarByClass(EmotionAnalysisDriver.class);
37         var4.setMapperClass(EmotionValidationMapper.class);
38         var4.setOutputKeyClass(Text.class);
39         var4.setOutputValueClass(Text.class);
40         FileInputFormat.addInputPath(var4, new Path(var0[1]));
41         FileOutputFormat.setOutputPath(var4, new Path(var0[2]));
42         var4.addCacheFile((new Path("/user/hadoop/emotion_training_output/part-r-00000")).toUri());
43         System.out.println("Starting validation phase...");
44         if (!var4.waitForCompletion(true)) {
45             System.err.println("Validation phase failed.");
46             System.exit(1);
47         }
48
49         System.out.println("Emotion Analysis completed successfully.");
50     }
51 }
```



# Chapter 7

## RESULT ANALYSIS IN HADOOP

This section provides an overview of the sentiment and emotion analysis operations performed using Hadoop, along with descriptions of the outputs generated. Screenshots showcase results at various stages, including raw outputs, trained outputs, and categorized data.

### 1. Sentiment Analysis Operation

#### Description:

The sentiment analysis operation processes input text data to classify sentiments into categories like Positive, Negative, and Neutral. It involves training a model to identify sentiment patterns and validating outputs against real-world data.

#### a. Sentiment Trained Output:

Highlights the refined sentiment results after the training process, showcasing improved accuracy in categorization and better distribution insights.

A screenshot of a terminal window displaying the output of a sentiment analysis operation. The output consists of a list of words followed by their assigned sentiment category (neutral, positive, or negative). The words are listed in a single column, and the sentiment categories are listed in a second column to the right of the words. The output is as follows:

wisely	neutral
wish	positive
wishes	neutral
wishing	neutral
wisma	positive
wisuda	positive
wit	neutral
witchu	negative
witchy	positive
with	positive
within	positive
without	neutral
witnessed	negative
witnessing	positive
wittaya	positive
wiv	neutral
wives	positive
wixy	negative

## b. Sentiment Output:

Displays the raw sentiment analysis results, showing the classification of text data into sentiment categories with their respective counts.

```
ilo hahaha @luanabilro positive
Quiero jugar tenis r0ñ positive
RT @KaroliyaAmeOub @CataDelina se los pegare en tu nombre en cualquier momento xD// Me parece perfectooooo xD jajajaj
neutral
Rylan... love you @bbuk xx positive
Safe to say I teared up more than I do in the average year in one night at Les Mis #musicgotme going #noshamer positive
Shopping! (@ Kohl's) http://t.co/I8ZkQHT9 neutral
Straight ignore. neutral
Time is a thing not to waste positive
Time to do work &gt;&gt;&gt;&gt;&gt; positive
Tirta gangga w/@dayu_renita @tamancrizz @mika_mahayekti http://t.co/xemAXjUz positive
Today I tripped and I'm 100% sure a fuck ton of people driving on Johnston saw it. My day&lt;&lt;&lt;&lt; positive
Tweet positive
Tweet 4484 Dedicado a @ZayraIzquierdo .* positive
Ufff que mal lo he paasado. #Muyfuertetodo positive
Ufff tenia que acabar mal el dia... positive
Ughh @HeavyHitterz91 needs to hurry the hell up and get outta work!! positive
Up =fAa neutral
Volkan konak adami tribe sokar yemin ederim :D positive
Vovozinha sua linda *----* http://t.co/RVWm594Q positive
Watching #Miranda On bbc1!!! @mermhart u r HILARIOUS fYh=f&i neutral
We rang 999 and they told us to go away? Rude that positive
Whistle baby neutral
Y para acabarla habla s||per lento mi compadre .... negative
Y tu el aire que me levanta que me da fuerza para este amor #TuAmorMeHaceTantoBien.. positive
YESSSSSSSSSS http://t.co/1qERLagY positive
Yeah fuck saquit fans i aint scared positive
Yessss ^ ^ positive
You never miss something til'ours gone positive
```

## 2. Emotion Analysis Operation

### Description:

The emotion analysis operation categorizes textual data into emotion labels such as Joy, Fear, or Disgust. The process includes training and validation steps to ensure accurate emotion classification.

- a. **Emotion Trained Output:** Displays the enhanced emotion analysis results post-training, demonstrating more precise emotion classification with aggregated summaries.

```
pong sadness
pongas fear
pongo anger
poniendo sadness
ponorogo surprise
pont fear
ponte joy
pontianak sadness
poo sadness
pool joy
pools surprise
poop sadness
pooper joy
poopos joy
poor fear
pootie joy
pop sadness
pope joy
popeyes anger
popo anger
popped disgust
pops surprise
popular joy
poquito joy
por joy
porc anger
porch sadness
pores anger
porfa joy
porfavor joy
```

## b. Emotion Output:

Contains the raw emotion classification results, showing the distribution of input data into predefined emotion categories.

```
Mampet hiding gw... surprise
Muy descontrolado sale Rolando Valdez #YaquisVsAlgodoneros surprise
My baby Is completely frozen no way I can drive this shit no time soon http://t.co/5fB1vytx joy
My ears once again hurt so much joy
M|d|l|r1|d-f|m|z|n 2013 y-1-1nda hayata ge|eirece-fi projeler |zerinde ki |eal-1fmalar sona yakla+f-1d-1.
joy
No no. I don't wanna get out of bed. /: gonna be a long dayy joy
Noooooossa!!! A garota acabou de lamber dentro da minha orelha!!! disgust
O.o neutral
Pequena no tamanho e cada vez maior nas suas conquistas. S|| orgulho!! Aguarde sabado a comemora|e|uo ser|i ao nosso est
ilo hahaha @_luanabilro disgust
Quiero jugar tenis r0N joy
RT @KaroliyaAmeOub @CataDeLina se los pegare en tu nombre en cualquier momento xD// Me parece perfectooooo xD jajajaj
sadness
Rylan... love you @bbuk xx joy
Safe to say I teared up more than I do in the average year in one night at Les Mis #musicgotmegoing #noshame joy
Shopping! (@ Kohl's) http://t.co/I8ZkQHT9 joy
Straight ignore. joy
Time is a thing not to waste joy
Time to do work &gt;&gt;&gt;&gt;&gt; joy
Tirta gangga w/@dayu_renita @tamancrizz @mika_mahayekti http://t.co/xemAXjUz joy
Today I tripped and I'm 100% sure a fuck ton of people driving on Johnston saw it. My day&lt;&lt;&lt; joy
Tweet sadness
Tweet 4484 Dedicado a @ZayraIzquierdo_* sadness
Ufff que mal lo he paasado. #Muyfuertetodo joy
Ufff tenia que acabar mal el dia... surprise
Ughh @HeavyHitterz91 needs to hurry the hell up and get outta work!! joy
Up =fAa neutral
Volkan konak adami tribe sokar yemin ederim :D fear
Vovozinha sua linda *----* http://t.co/RVWm594Q surprise
```

The results of the sentiment and emotion analysis operations demonstrate the effectiveness of the Hadoop framework for processing large-scale text data. Both operations were performed with a combination of mappers, reducers, and driver classes, which allowed for efficient data transformation, classification, and aggregation.

## Key Observations:

### 1. Sentiment Analysis:

- The initial sentiment output categorized input text into Positive, Negative, and Neutral sentiments, revealing a balanced distribution of sentiments in the dataset.
- The trained sentiment output provided refined classifications with improved accuracy, highlighting subtle variations in sentiment patterns that were less evident in raw data.
- The sentiment analysis process effectively identified key emotional trends in the input, aiding in understanding user feedback or public opinions.

### 2. Emotion Analysis:

- The raw emotion output classified textual data into emotion labels like Joy, Fear, and Disgust. The most dominant emotions were clearly visible, with some variations across different regions or contexts.
- The trained emotion output enhanced the precision of classifications by reducing ambiguity in emotion labeling. It provided a more accurate reflection of the emotional tone in the dataset, which is crucial for applications like customer experience analysis or social media monitoring.

## **Chapter 8**

### **CONCLUSION**

The sentiment analysis project utilizing Hadoop's MapReduce framework demonstrates an efficient and scalable approach to processing and analyzing vast volumes of Twitter data. By leveraging distributed computing, the system handles the challenges of large-scale data, providing accurate sentiment classification based on preprocessed datasets. The implementation ensures robustness and scalability, making it suitable for real-time or near-real-time analysis. This work highlights the effectiveness of big data technologies in transforming unstructured social media data into actionable insights, enabling informed decision-making for various applications like marketing, crisis management, and trend analysis.

### **FUTURE SCOPE**

While the current implementation focuses on sentiment classification using MapReduce, future advancements could include the integration of advanced machine learning techniques, such as deep learning, to enhance sentiment detection accuracy. Expanding the system to support multilingual analysis and real-time processing can broaden its applicability. Additionally, incorporating dynamic data visualization tools and dashboards can provide stakeholders with real-time sentiment trends. Adopting hybrid big data frameworks, such as Hadoop with Spark, could further improve processing speed and flexibility, paving the way for more complex and versatile analytics pipelines.

# Chapter 9

## REFERENCE

- [1] Dean, J., & Ghemawat, S. (2008). MapReduce: Simplified data processing on large clusters. Communications of the ACM, 51(1), 107–113. <https://doi.org/10.1145/1327452.1327492>
- [2] White, T. (2015). Hadoop: The definitive guide. O'Reilly Media, Inc.
- [3] Lin, J., & Dyer, C. (2010). Data-intensive text processing with MapReduce. Morgan & Claypool Publishers.
- [4] Borthakur, D. (2008). The Hadoop distributed file system: Architecture and design. Apache Software Foundation. Retrieved from <https://hadoop.apache.org>
- [5] Zaharia, M., Chowdhury, M., Franklin, M. J., Shenker, S., & Stoica, I. (2010). Spark: Cluster computing with working sets. Proceedings of the 2nd USENIX Conference on Hot Topics in Cloud Computing (HotCloud '10).
- [6] Aggarwal, C. C., & Zhai, C. (2012). Mining text data. Springer Science & Business Media. <https://doi.org/10.1007/978-1-4614-3223-4>
- [7] Al-Msie'deen, R., Awajan, A., & Moghrabi, A. (2020). Sentiment analysis and emotion detection in social media text. Journal of Information Science, 46(3), 383–399. <https://doi.org/10.1177/0165551519837309>
- [8] Cambria, E., Schuller, B., Xia, Y., & Havasi, C. (2013). New avenues in opinion mining and sentiment analysis. IEEE Intelligent Systems, 28(2), 15–21. <https://doi.org/10.1109/MIS.2013.30>
- [9] Liu, B. (2012). Sentiment analysis and opinion mining. Synthesis Lectures on Human Language Technologies, 5(1), 1–167. <https://doi.org/10.2200/S00416ED1V01Y201204HLT016>
- [10] Mikolov, T., Sutskever, I., Chen, K., Corrado, G. S., & Dean, J. (2013). Distributed representations of words and phrases and their compositionality. Advances in Neural Information Processing Systems, 26, 3111–3119.